sealed

Testing an AI safely — explained in one page.

The problem

To test an AI, one team tries to trick it and another team defends it — at the same time.

Why that ruins the test

- The defenders can see the attack before the test ends — so a “pass” is fake.
- The AI can carry memory from one test into the next and get confused.
- Sharing one machine opens subtle leaks that normal setups ignore.

The bottom line: if the teams can peek at each other, the safety result cannot be trusted.

The solution

Keep each team in its own room, and route everything through one trusted go-between.

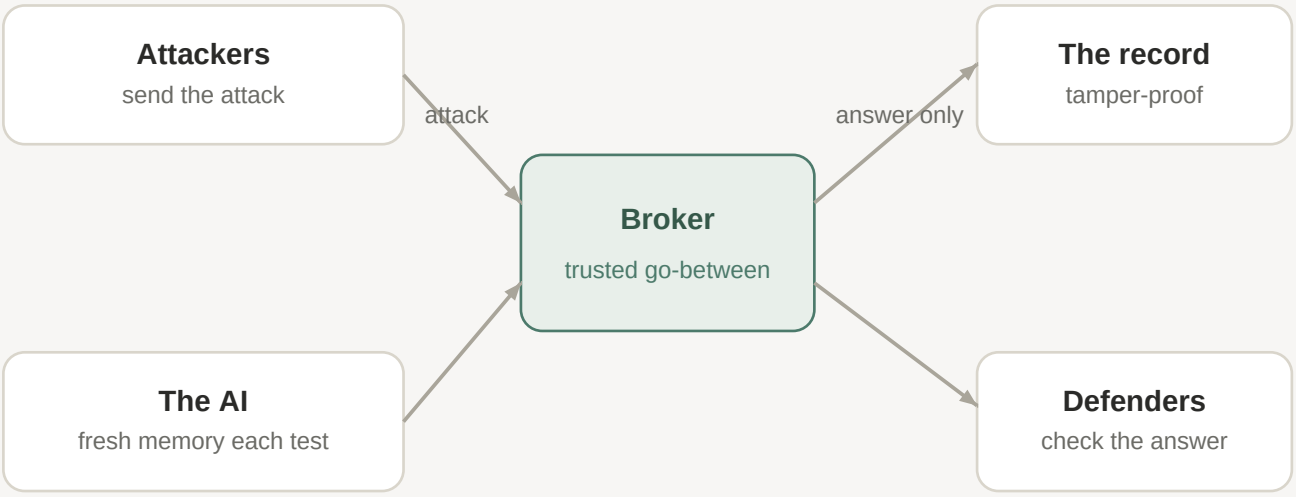

How it keeps the test honest

- Each team works in its own locked room — they never talk to each other.
- They only talk to the go-between, which decides who may see what.
- The defenders get only the answer — never the raw attack — until the test ends.
- The AI starts every test with a blank memory, so nothing carries over.
- Every step is saved in a record that can't be secretly changed.

It stops

- teams peeking at each other
- one test leaking into the next
- anyone editing the record
- unauthorized requests

It doesn't claim

- deep chip-level leaks
- a hacked server
- a dishonest go-between

Being honest about the limits is what makes the rest trustworthy.